



```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
```

```
In [2]: movies = pd.read_csv('movie.csv')
```

```
In [3]: movies.head()
```

```
Out[3]:
```

	movield	title	genres
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2	Jumanji (1995)	Adventure Children Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama Romance
4	5	Father of the Bride Part II (1995)	Comedy

```
In [4]: movies.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27278 entries, 0 to 27277
Data columns (total 3 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   movieId  27278 non-null   int64  
 1   title    27278 non-null   object  
 2   genres   27278 non-null   object  
dtypes: int64(1), object(2)
memory usage: 639.5+ KB
```

```
In [5]: ratings = pd.read_csv('rating.csv')
```

```
In [6]: ratings.head()
```

```
Out[6]:
```

	userId	movield	rating	timestamp
0	1	2	3.5	2005-04-02 23:53:47
1	1	29	3.5	2005-04-02 23:31:16
2	1	32	3.5	2005-04-02 23:33:39
3	1	47	3.5	2005-04-02 23:32:07
4	1	50	3.5	2005-04-02 23:29:40

```
In [7]: ratings.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000263 entries, 0 to 20000262
Data columns (total 4 columns):
 #   Column      Dtype  
--- 
 0   userId      int64  
 1   movieId     int64  
 2   rating      float64 
 3   timestamp   object  
dtypes: float64(1), int64(2), object(1)
memory usage: 610.4+ MB
```

```
In [8]: df = ratings.merge(movies, on="movieId", how="left")
df.head()
```

	userId	movieId	rating	timestamp	title	gen
0	1	2	3.5	2005-04-02 23:53:47	Jumanji (1995)	Adventure Children Fant
1	1	29	3.5	2005-04-02 23:31:16	City of Lost Children, The (Cité des enfants p...)	Adventure Drama Fantasy Mystery
2	1	32	3.5	2005-04-02 23:33:39	Twelve Monkeys (a.k.a. 12 Monkeys) (1995)	Mystery Sci-Fi Thr
3	1	47	3.5	2005-04-02 23:32:07	Seven (a.k.a. Se7en) (1995)	Mystery Thr
4	1	50	3.5	2005-04-02 23:29:40	Usual Suspects, The (1995)	Crime Mystery Thr

```
In [9]: import re
df['year'] = df['title'].str.extract(r'\((\d{4})\)').astype(float)
df.head()
```

Out[9]:

	userId	movieId	rating	timestamp	title	gen
0	1	2	3.5	2005-04-02 23:53:47	Jumanji (1995)	Adventure Children Fant
1	1	29	3.5	2005-04-02 23:31:16	City of Lost Children, The (Cité des enfants p...)	Adventure Drama Fantasy Mystery
2	1	32	3.5	2005-04-02 23:33:39	Twelve Monkeys (a.k.a. 12 Monkeys) (1995)	Mystery Sci-Fi Thr
3	1	47	3.5	2005-04-02 23:32:07	Seven (a.k.a. Se7en) (1995)	Mystery Thr
4	1	50	3.5	2005-04-02 23:29:40	Usual Suspects, The (1995)	Crime Mystery Thr

```
In [10]: df['genres'] = df['genres'].str.split('|')
df.head()
```

Out[10]:

	userId	movieId	rating	timestamp	title	genres	year
0	1	2	3.5	2005-04-02 23:53:47	Jumanji (1995)	[Adventure, Children, Fantasy]	1995.0
1	1	29	3.5	2005-04-02 23:31:16	City of Lost Children, The (Cité des enfants p...)	[Adventure, Drama, Fantasy, Mystery, Sci-Fi]	1995.0
2	1	32	3.5	2005-04-02 23:33:39	Twelve Monkeys (a.k.a. 12 Monkeys) (1995)	[Mystery, Sci-Fi, Thriller]	1995.0
3	1	47	3.5	2005-04-02 23:32:07	Seven (a.k.a. Se7en) (1995)	[Mystery, Thriller]	1995.0
4	1	50	3.5	2005-04-02 23:29:40	Usual Suspects, The (1995)	[Crime, Mystery, Thriller]	1995.0

```
In [11]: df.rename(columns={'timestamp': 'date'}, inplace=True)
df.head()
```

Out[11]:

	userId	movieId	rating	date	title	genres	year
0	1	2	3.5	2005-04-02 23:53:47	Jumanji (1995)	[Adventure, Children, Fantasy]	1995.0
1	1	29	3.5	2005-04-02 23:31:16	City of Lost Children, The (Cité des enfants p...)	[Adventure, Drama, Fantasy, Mystery, Sci-Fi]	1995.0
2	1	32	3.5	2005-04-02 23:33:39	Twelve Monkeys (a.k.a. 12 Monkeys) (1995)	[Mystery, Sci- Fi, Thriller]	1995.0
3	1	47	3.5	2005-04-02 23:32:07	Seven (a.k.a. Se7en) (1995)	[Mystery, Thriller]	1995.0
4	1	50	3.5	2005-04-02 23:29:40	Usual Suspects, The (1995)	[Crime, Mystery, Thriller]	1995.0

In [12]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000263 entries, 0 to 20000262
Data columns (total 7 columns):
 #   Column    Dtype  
 --- 
 0   userId     int64  
 1   movieId    int64  
 2   rating     float64 
 3   date       object  
 4   title      object  
 5   genres     object  
 6   year       float64 
dtypes: float64(2), int64(2), object(3)
memory usage: 1.0+ GB
```

In [13]: `df.isnull().sum()`


```
Out[13]: userId      0
          movieId     0
          rating      0
          date        0
          title       0
          genres      0
          year        405
          dtype: int64
```

In [16]: `df['year'].fillna(df['year'].mode()[0])`
`df.isnull().sum()`

```
Out[16]: userId      0
          movieId     0
          rating      0
          date        0
          title       0
          genres      0
          year        0
          dtype: int64
```

```
In [15]: movies_df = df.copy() #I am Using .copy() for protects the ori
```

```
In [16]: plt.figure(figsize=(8,5))
sns.histplot(movies_df['rating'], bins=20, kde=True)
plt.title('Distribution of User Ratings')
plt.xlabel('Rating')
plt.ylabel('count')
plt.show()
```



```
In [17]: # The distribution shows that most movie ratings fall between 3.0 and 4.0,
# indicating that users tend to rate movies favorably.
# Extreme ratings (very low or very high) are relatively rare.
# The scientific notation 1e7 on the y-axis represents counts in the scale of
# meaning the dataset contains a large number of total ratings.
```

```
In [18]: # The dataset is too large, so I am working with a sample for genre analysis.
# I am running into a MemoryError because exploding the genres column creates
```

```
sample_df = movies_df.sample(400_000, random_state=42)
movies_genres = sample_df[['movieId', 'rating', 'genres']].explode('genres')


movies_genres.head()
```

Out[18]:

	movielid	rating	genres
17679788	8360	3.5	Adventure
17679788	8360	3.5	Animation
17679788	8360	3.5	Children
17679788	8360	3.5	Comedy
17679788	8360	3.5	Musical

In [19]:

```
plt.figure(figsize=(12,5))
movies_genres['genres'].value_counts().head(10).plot(kind="bar", color = 'red')
plt.title("Top 10 Most Common Genres")
plt.ylabel("Count")
plt.show()
```




In [20]:

```
genre_ratings = movies_genres.groupby('genres')['rating'].mean().sort_values(ascending=False)
genre_ratings.head(10)
```

```
Out[20]: genres
Film-Noir      3.983699
War           3.815908
Documentary   3.764737
Crime          3.679582
Drama          3.676948
Mystery         3.669854
IMAX           3.665881
Animation       3.616760
Western         3.563012
Musical          3.560036
Name: rating, dtype: float64
```

```
In [21]: plt.figure(figsize=(12,5))
genre_ratings.head(10).plot(kind='bar', color='green')
plt.title("Top 10 Highest Rated Genres")
plt.ylabel("Average Rating")
plt.show()
```



```
In [22]: movie_stats = sample_df.groupby('title').agg({
    'rating': 'mean',
    'userId': 'count'
}).rename(columns={'userId': 'rating_count'})

filtered_movies = movie_stats[movie_stats['rating_count'] >= 50]

top_movies = filtered_movies.sort_values(by='rating', ascending=False).head(10)
top_movies
```

Out[22]:

		rating	rating_count
	title		
	Shawshank Redemption, The (1994)	4.496828	1261
	Usual Suspects, The (1995)	4.378885	933
	Double Indemnity (1944)	4.357759	116
	Godfather, The (1972)	4.350059	847
	Schindler's List (1993)	4.315738	1023
	Sunset Blvd. (a.k.a. Sunset Boulevard) (1950)	4.308824	136
	Godfather: Part II, The (1974)	4.308608	546
	Lives of Others, The (Das leben der Anderen) (2006)	4.297414	116
	Big Sleep, The (1946)	4.295918	98
	Rear Window (1954)	4.293716	366

In [23]: `# Groups movies → calculates their average rating and number of ratings → removes them by rating → shows the top 10 highest-rated movies.`

In [24]: `sample_df.head()`

Out[24]:

	userId	movieId	rating	date	title	genres	year
	17679788	122270	8360	3.5 2012-04-22 01:07:04	Shrek 2 (2004)	[Adventure, Animation, Children, Comedy, Music...]	2004.0
	7106385	49018	32	2.0 2001-09-11 07:50:36	Twelve Monkeys (a.k.a. 12 Monkeys) (1995)	[Mystery, Sci-Fi, Thriller]	1995.0
	12970708	89527	109374	3.5 2015-01-06 09:26:40	Grand Budapest Hotel, The (2014)	[Comedy, Drama]	2014.0
	15426752	106704	1060	3.0 2000-01-22 21:27:57	Swingers (1996)	[Comedy, Drama]	1996.0
	6934678	47791	1732	2.0 2006-01-19 15:48:23	Big Lebowski, The (1998)	[Comedy, Crime]	1998.0


In [25]: `rating_trend = sample_df.groupby('year')['rating'].mean().reset_index()`

```
rating_trend.head()
```

```
Out[25]:      year    rating
0  1896.0  3.000000
1  1902.0  4.333333
2  1903.0  3.750000
3  1909.0  3.000000
4  1914.0  2.750000
```


```
In [26]: rating_trend.plot(x='year', y='rating', kind='line', figsize=(10,5), title='Tr
```

```
Out[26]: <Axes: title={'center': 'Trend of Average Movie Ratings Over Years'}, xlabel='year'>
```



```
In [27]: popularity_trend = sample_df.groupby('year')['userId'].count().reset_index()
popularity_trend.plot(x='year', y='userId', kind='bar', figsize=(12,5), title=
```


```
Out[27]: <Axes: title={'center': 'Number of Ratings Per Year (Movie Popularity Trend)'}, xlabel='year'>
```



```
In [28]: genre_year_trend = sample_df.explode('genres').groupby(['year', 'genres'])['rating'].mean().reset_index()
genre_year_trend.head(10)
```

	year	genres	rating
0	1896.0	Documentary	3.000000
1	1902.0	Action	4.333333
2	1902.0	Adventure	4.333333
3	1902.0	Fantasy	4.333333
4	1902.0	Sci-Fi	4.333333
5	1903.0	Crime	3.750000
6	1903.0	Western	3.750000
7	1909.0	Crime	3.000000
8	1909.0	Drama	3.000000
9	1914.0	Comedy	2.500000

```
In [29]: plt.figure(figsize=(12,6))
sns.lineplot(data=genre_year_trend, x='year', y='rating', hue='genres')
plt.title("Genre Rating Trends Over Years")
plt.legend(bbox_to_anchor=(1,1))
plt.show()
```



```
In [30]: movie_year_stats = sample_df.groupby(['title', 'year'])['userId'].count().reset_index()
top_movies = movie_year_stats.sort_values(by='userId', ascending=False).head(10)
top_movies
```

Out[30]:

	title	year	userId
8879	Pulp Fiction (1994)	1994.0	1389
4081	Forrest Gump (1994)	1994.0	1287
9956	Silence of the Lambs, The (1991)	1991.0	1264
9849	Shawshank Redemption, The (1994)	1994.0	1261
5991	Jurassic Park (1993)	1993.0	1175
10368	Star Wars: Episode IV - A New Hope (1977)	1977.0	1114
10852	Terminator 2: Judgment Day (1991)	1991.0	1107
7132	Matrix, The (1999)	1999.0	1086
1651	Braveheart (1995)	1995.0	1030
9630	Schindler's List (1993)	1993.0	1023

```
In [31]: sample_df.head()
```

Out[31]:

	userId	movieId	rating	date	title	genres	year
17679788	122270	8360	3.5	2012-04-22 01:07:04	Shrek 2 (2004)	[Adventure, Animation, Children, Comedy, Music...]	2004.0
7106385	49018	32	2.0	2001-09-11 07:50:36	Twelve Monkeys (a.k.a. 12 Monkeys) (1995)	[Mystery, Sci-Fi, Thriller]	1995.0
12970708	89527	109374	3.5	2015-01-06 09:26:40	Grand Budapest Hotel, The (2014)	[Comedy, Drama]	2014.0
15426752	106704	1060	3.0	2000-01-22 21:27:57	Swingers (1996)	[Comedy, Drama]	1996.0
6934678	47791	1732	2.0	2006-01-19 15:48:23	Big Lebowski, The (1998)	[Comedy, Crime]	1998.0

In [49]: `mr_df = sample_df[['userId', 'movieId', 'rating']].copy()`

In [50]: `mr_df['userId'] = mr_df['userId'].astype('category').cat.codes
mr_df['movieId'] = mr_df['movieId'].astype('category').cat.codes
mr_df.head()`

Out[50]:

	userId	movieId	rating
17679788	87938	6809	3.5
7106385	35270	31	2.0
12970708	64366	12084	3.5
15426752	76718	981	3.0
6934678	34378	1558	2.0

In [51]: `from sklearn.model_selection import train_test_split

x = mr_df[['userId', 'movieId']]
y = mr_df['rating']

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)`

In [52]: `from sklearn.linear_model import LinearRegression`

```
lm = LinearRegression()  
lm.fit(x_train, y_train)
```

Out[52]:

LinearRegression		
Parameters		
fit_intercept	True	
copy_X	True	
tol	1e-06	
n_jobs	None	
positive	False	

In [53]: `y_pred = lm.predict(x_test)`

In [54]: `from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score`

```
rmse = np.sqrt(mean_squared_error(y_test, y_pred))  
mae = mean_absolute_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)  
  
print("Model Evaluation Results")  
print("RMSE:", rmse)  
print("MAE:", mae)  
print("R2 Score:", r2)
```

Model Evaluation Results
RMSE: 1.0540940457331724
MAE: 0.8420110436343979
R2 Score: 0.00024800373574662693

In [55]: `from sklearn.linear_model import Ridge`

```
ridge_model = Ridge(alpha=1.0)  
ridge_model.fit(x_train, y_train)  
  
ridge_pred = ridge_model.predict(x_test)  
  
rmse_ridge = np.sqrt(mean_squared_error(y_test, ridge_pred))  
print("Ridge RMSE:", rmse_ridge)
```

Ridge RMSE: 1.0540940457331727

In [56]: `from sklearn.linear_model import Lasso`

```
lasso_model = Lasso(alpha=0.001)  
lasso_model.fit(x_train, y_train)
```

```
lasso_pred = lasso_model.predict(x_test)

rmse_lasso = np.sqrt(mean_squared_error(y_test, lasso_pred))
print("Lasso RMSE:", rmse_lasso)
```

```
Lasso RMSE: 1.0540940473418565
```

```
In [57]: models = {
    "Linear Regression": rmse,
    "Ridge Regression": rmse_ridge,
    "Lasso Regression": rmse_lasso
}

models
```

```
Out[57]: {'Linear Regression': 1.0540940457331724,
'Ridge Regression': 1.0540940457331727,
'Lasso Regression': 1.0540940473418565}
```

```
In [58]: #All three models show almost identical performance.
#This indicates the dataset is simple, the features are not highly correlated,
# let's upgrade movie rating prediction model with feature engineering + better
```

```
In [59]: # Feature Engineering

# Average rating per movie
sample_df['movie_avg_rating'] = sample_df.groupby('movieId')['rating'].transform('mean')

# Average rating per user
sample_df['user_avg_rating'] = sample_df.groupby('userId')['rating'].transform('mean')

# Rating count per movie
sample_df['movie_rating_count'] = sample_df.groupby('movieId')['rating'].transform('count')

# Rating count per user
sample_df['user_rating_count'] = sample_df.groupby('userId')['rating'].transform('count')
```

```
In [60]: from sklearn.model_selection import train_test_split

features = [
    'userId', 'movieId',
    'movie_avg_rating', 'user_avg_rating',
    'movie_rating_count', 'user_rating_count'
]

X = sample_df[features]
y = sample_df['rating']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [61]: from sklearn.preprocessing import OrdinalEncoder

encoder = OrdinalEncoder()
```

```
    handle_unknown='use_encoded_value',
    unknown_value=-1
)

X_train.loc[:, ['userId','movieId']] = encoder.fit_transform(
    X_train[['userId','movieId']])
)

X_test.loc[:, ['userId','movieId']] = encoder.transform(
    X_test[['userId','movieId']])
)
```

```
In [62]: from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error
import numpy as np

models = {
    "Linear Regression": LinearRegression(),
    "Ridge Regression": Ridge(alpha=1.0),
    "Lasso Regression": Lasso(alpha=0.001)
}

results = {}

for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    rmse = np.sqrt(mean_squared_error(y_test, preds))
    results[name] = rmse
```

```
In [63]: results
```

```
Out[63]: {'Linear Regression': 0.7546256717824902,
          'Ridge Regression': 0.7546256943450338,
          'Lasso Regression': 0.7546394173818821}
```